AutoTriager

Evidence-First Incident Triage for Microservices

Dahe Chen · University of Missouri–Kansas City · CS 5588 Capstone · Hack-A-Roo

CS 5588 - Data Science Capstone | Instructor: Yugyung Lee

735
Failure cases in the full RCAEval benchmark

90
RE2 Train Ticket cases in the published study

6
Fault types in the published Train Ticket comparison

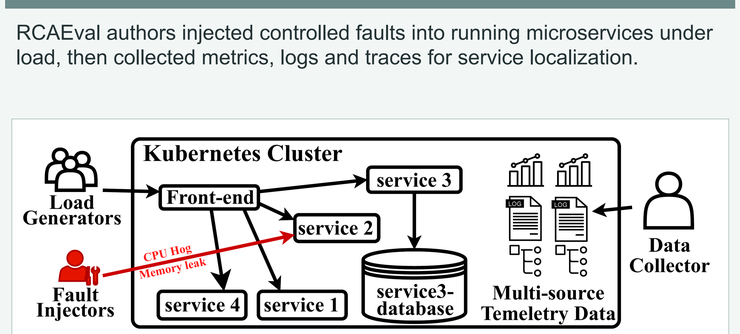
01 Problem

Which service should the engineer check first?

A slow storefront can reflect faults in different downstream services. AutoTriager gives the on-call engineer a ranked shortlist and the observations behind it.

02 Benchmark Data

RCAEval authors injected controlled faults into running microservices under load, then collected metrics, logs and traces for service localization.



RCAEval data collection setup. Pham et al. [1], Figure 2. Cropped from the original, CC BY 4.0.

03 Published Evaluation Protocol

Cohort

Pham et al. compare methods on RE2 Train Ticket. The 90 cases cover six fault types, five services and three repeats per combination.

Ranking metrics

AC@1 measures how often the injected service ranks first. AC@3 measures top-three coverage. Avg@5 averages AC@1 through AC@5.

Compared methods

The paper compares metric-based, trace-based and multi-source baselines for ranking fault services. The chart selects two published methods.

Source: RCAEval [1], Table 2, §4.2 and §5.

04 AutoTriager B1 Workflow

01 Bind the incident window

Scope available telemetry to the incident window. Keep service and source identities.

↓

02 Assemble fixed evidence

Keep up to 12 standardized metric shifts, 8 log groups and 8 observed call links.

↓

03 Rank candidates in one pass

Gemini 3.5 Flash-Lite reads the complete packet and returns up to five distinct services.

↓

04 Validate and return

Check the schema, service names and evidence IDs. Present the ranked services for review.

Output Ranked services, evidence links and suggested checks

AutoTriager implements the fixed-packet workflow. The published baseline study shown here was conducted by the RCAEval authors.

05 Workbench Demonstration

Open the incident

Review the affected service and the available observation window.

Compare ranked services

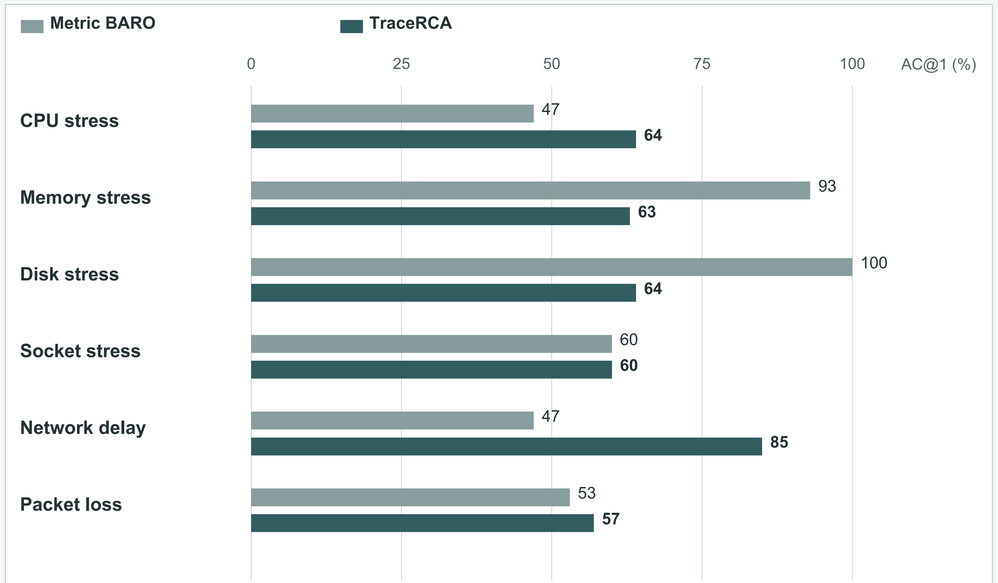
Read the saved ranking and the observations cited for each candidate.

Inspect the evidence

Check the underlying observations before deciding which service to investigate.

Demonstration uses a recorded case.

06 Published Benchmark Results · RE2 Train Ticket



Published first-place service accuracy by fault type. Redrawn from Pham et al. [1], Table 6, CC BY 4.0.

On disk, BARO scored 100% versus 64%. On delay, TraceRCA scored 85% versus 47%.

07 Published Averages · Selected Baselines

The table reports the paper's average columns for RE2 Train Ticket.

Input	Method	AC@1	AC@3	Avg@5
Metric	BARO	0.67	0.82	0.80
Trace	TraceRCA	0.66	0.79	0.77
Multi-source	BARO	0.69	0.82	0.81

Redrawn from Pham et al. [1], Table 6, average columns, CC BY 4.0.

The chart and table report the RCAEval authors' benchmark experiment. AutoTriager B1 is presented separately as the workbench's investigation method.

Published evidence shows that service localization performance varies with fault type.

[1] Pham et al. RCAEval. WWW Companion 2025. DOI 10.1145/3701716.3715290. Original Figure 2 and both redraws of Table 6 are attributed under CC BY 4.0.